

How to select predictive models for decision-making or causal inference?

Matthieu Doutreligne^{1,2,*} and Gaël Varoquaux¹

¹Inria, Soda, Saclay, France

²Mission Data, Haute Autorité de Santé, Saint-Denis, France

*Corresponding author: matt.dout@gmail.com

December 8, 2023

Abstract

Objective: As predictive models –eg from machine learning– give likely outcomes, they may be used to reason on the effect of an intervention, a causal-inference task. The increasing complexity of health data has opened the door to a plethora of models, but also the Pandora box of model selection: which of these models yield the most valid causal estimates? In this paper, we study model selection procedures in practical settings: finite samples settings and without well-specification assumption. **Methods:** Elaborate risks build proxies of the causal error using “nuisance” re-weighting to compute it on the observed data. But does computing these nuisance adds noise to model selection? Drawing from an extensive empirical study based on a simulation and three semi-simulated healthcare datasets based on real covariates, we compare different metrics for causal model selection.

Results: Among all metrics the classical mean squared error (ie. factual μ -risk) is worse and re-weighting it with propensity score (μ -risk_{IPW}) does not bring much improvements. The R -risk, which includes a model of mean outcome and propensity scores, leads to the best performances.

Conclusions: We outline a good causal model-selection procedure using the so-called R -risk; using flexible estimators to compute the nuisance models on the train set; and splitting out 10% of the data to compute risks.

Keywords Model Selection, Treatment Effect, G-formula, Observational Study, Machine Learning,

Background

Extending prediction to prescription needs causality

Progress in machine learning brings predictive models to new health data [10, 63], with the promise of precision medicine. Predictive models are increasingly used for image analysis [39, 94, 92, 76, 52], routinely-collected patient records [49, 19, 80] or clinical notes [32, 90, 82]. Data may be difficult to control and model, but the accuracy of the prediction can be verified on left-out data [3, 60, 88]. Given a model predicting a health outcome, precision medicine would like it to guide decisions: will an individual benefit from an intervention such as surgery [22]? Contrasting predictions with and without the treatment gives an answer, but statistical validity requires causal inference [81, 83, 11].

Predictive modeling bridges to causal-inference with the so-called *outcome models*, or G-computation, G-formula [65], Q-model [81], conditional mean regression [91]. A central challenge of inference of treatment effects is that of confounding: spurious associations between treatment allocation and baseline health, *eg* only prescribing a drug to mild cases [29, 87]. Causal inference has been central to epidemiology, typically with methods that model treatment assignment [8, 13, 25], based on propensity scores [68]. But outcome modeling go beyond average effects, estimating individualized or conditional average treatment effects (CATE), central to precision medicine. For this purpose, such methods are also invaluable on randomized trials [85, 44, 31].

Outcome-modeling methods, even when specifically designed for causal inference are numerous: Bayesian Additive Regression Trees [30], Targeted Maximum Likelihood Estimation [43, 74], causal boosting [61], causal multivariate adaptive regression splines [61], random forests [89, 5], Meta-learners [40], R-learners [54], Doubly robust estimation [15]... The wide variety of methods raises the problem of selecting between different estimators based on the data at hand. Indeed, estimates of treatment effects can vary markedly across different predictive models [21, 20] (illustration in Appendix A).

Given complex, potentially noisy, data, which model is to be most trusted to yield valid causal estimates? As no single learner performs best on all data sets, there is a pressing need for clear guidelines to select outcome models for causal inference.

Standard practices to select models in predictive settings rely on the error on the outcome [60, 88]. However, as we will see, these practices may not pick the best models for causal inference, as they can be misled by inhomogeneities due to treatment allocation.

Objectives and structure of the paper We study *model selection procedures* in practical settings: *finite samples* settings and without *well-specification* assumption. Asymptotic causal-inference theory calls for complex risks, but a practical question is whether model-selection procedures, that rely on data split, can estimate these risks reliably enough. Indeed,

they come with more quantities to estimate, which may bring additional variance, leading to worse model selection.

We first illustrate the problem of causal model selection and briefly review prior art. Then we anchor causal model selection in the *potential outcome* framework and details the causal risks and model-selection procedure. We then rewrite the so-called *R-risk*, as a reweighted version of mean squared difference between the true and estimated individualized treatment effect. Finally, we conduct a thorough empirical study comparing the different metrics on a wide variety of data-generative processes based on simulations and real data.




Figure 1: **Illustration:** a) a random-forest estimator with high performance for standard prediction (high R^2) but that yields poor causal estimates (large error between true effect τ and estimated $\hat{\tau}$), b) a linear estimator with smaller prediction performance leading to better causal estimation.

Selecting the estimator with the smallest error to the individual treatment effect $\mathbb{E}[(\tau(x) - \hat{\tau}(x))^2]$ – the τ -risk, eq. 9 – would lead to the best causal estimates; however computing this error is not feasible: it requires access to unknown quantities: $\tau(x)$.

While the random forest fits the data better than the linear model, it gives worse causal inference because its error is inhomogeneous between treated and untreated. The R^2 score does not capture this inhomogeneity.

Illustration: the best predictor may not estimate best causal effects

Using a predictor to reason on causal effects relies on contrasting the prediction of the outcome for a given individual with and without the treatment. Given various predictors of the outcome, which one should we use? Standard predictive modeling or machine-learning practice selects the predictor that minimizes the expected error. However, this predictor may not be the best model to reason about causal effects of an intervention as Figure 1 illustrates. Consider the probability Y of an undesirable outcome (eg death), a binary treatment $A \in \{0, 1\}$, and a covariate $X \in \mathbb{R}$ summarizing the patient health status (eg. the Charlson index [14]). We simulate a treatment beneficial (decreases mortality) for patients with high Charlson scores (bad health status) but with little effect for patients in good condition (low Charlson scores).

Figure 1a shows a random forest predictor with a counter-intuitive behavior: it predicts well on average the outcome (as measured by a regression R^2 score) but perform poorly to estimate causal quantities: the average treatment effect τ (as visible via the error $|\tau - \hat{\tau}|$) or the conditional average treatment effect (the error $\mathbb{E}[(\tau(x) - \hat{\tau}(x))^2]$, called CATE). On the contrary, Figure 1b shows a linear model with smaller R^2 score but better causal inference.

The problem is that causal estimation requires controlling an error on both treated and non-treated outcome for the same individual: the observed outcome, and the non-observed *counterfactual* one. The linear model is misspecified –the outcome functions are not linear–, leading to poor R^2 ; but it interpolates better to regions where there are few untreated individuals –high Charlson score– and thus gives better causal estimates. Conversely, the random forest puts weaker assumptions on the data, thus has higher R^2 score but is biased by the treated population in the poor-overlap region, leading to bad causal estimates.

This toy example illustrates that the classic minimum Mean Square Error criterion is not suited to choosing a model among candidate estimators for causal inference.

Prior work: model selection for outcome modeling (g-computation)

A natural way to select a predictive model for causal inference would be an error measure between a causal quantity such as the CATE and models' estimate. But such error is not a “feasible” risk: it cannot be computed solely from observed data and requires oracle knowledge.

Simulation studies of causal model selection Using eight simulations setups from [61], where the oracle CATE is known, Schuler et al. [73] compare four causal risks, concluding that for CATE estimation the best model-selection risk is the so-called R -risk [54] –def. 5, below. Their empirical results are clear for randomized treatment allocation but less convincing for observational settings where both simple Mean

Squared Error –MSE, μ -risk(f) def. 1– and reweighted MSE – μ -risk_{IPW} def. 2– appear to perform better than R -risk on half of the simulations. Another work [1] studied empirically both MSE and reweighted MSE risks on the semi-synthetic ACIC 2016 datasets [20], but did not include the R -risk. We complete these prior empirical work by studying a wider variety of data generative processes and varying the influence of overlap, an important parameter of the data generation process which makes a given causal metric appropriate [18]. We also study how to best adapt cross-validation procedures to causal metrics which themselves come with models to estimate.

Theoretical studies of causal model selection Several theoretical works have proposed causal model selection procedures that are *consistent*: select the best model in a family given asymptotically large data. These work rely on introducing a CATE estimator in the testing procedure: matching [67], an IPW estimate [27], a doubly robust estimator [71], or debiasing the error with influence functions [1]. However, for theoretical guarantees to hold, the test-set correction needs to converge to the oracle: it needs to be flexible enough –well-posed– and asymptotic data. From a practical perspective, meeting such requirements implies having a good CATE estimate, thus having solved the original problem of causal model selection.

Statistical guarantees on causal estimation procedures Much work in causal inference has focused on procedures that guarantee asymptotically consistent estimators, such as Targeted Machine Learning Estimation (TMLE) [43, 74] or Double Machine Learning [15]. Here also, theories require asymptotic regimes and models to be *well-specified*.

By contrast, Johansson et al. [37] studies causal estimation without assuming that estimators are well specified. They derive an upper bound on the oracle error to the CATE (τ -risk) that involves the error on the outcome and the similarity of the distributions of treated and control patients. However, they use this upper bound for model optimization, and do not give insights on model selection. In addition, for hyperparameter selection, they rely on a plugin estimate of the τ -risk built with counterfactual nearest neighbors, which has been shown ineffective [73].

Methods

Neyman-Rubin Potential Outcomes framework

Settings The Neyman-Rubin Potential Outcomes framework [51, 34] enables statistical reasoning on causal treatment effects: Given an outcome $Y \in \mathbb{R}$ (eg. mortality risk or hospitalization length), function of a binary treatment $A \in \mathcal{A} = \{0, 1\}$ (eg. a medical procedure), and baseline covariates $X \in \mathcal{X} \subset \mathbb{R}^d$, we observe the factual distribution, $O = (Y(A), X, A) \sim \mathcal{D} = \mathbb{P}(y, x, a)$. However, we want to model the existence of potential observations (unobserved ie. counterfactual) that correspond to a different treatment.

Thus we want quantities on the counterfactual distribution $O^* = (Y(1), Y(0), X, A) \sim \mathcal{D}^* = \mathbb{P}(y(1), y(0), x, a)$.

Popular quantities of interest (estimands) are: at the population level, the Average Treatment Effect

$$\text{ATE} \quad \tau \stackrel{\text{def}}{=} \mathbb{E}_{Y(1), Y(0) \sim \mathcal{D}^*}[Y(1) - Y(0)];$$

at the individual level, to model heterogeneity, the Conditional Average Treatment Effect

$$\text{CATE} \quad \tau(x) \stackrel{\text{def}}{=} \mathbb{E}_{Y(1), Y(0) \sim \mathcal{D}^*}[Y(1) - Y(0)|X = x].$$

Causal assumptions Assumptions are necessary for causal estimands to be identifiability in observational settings [70]. We assume the usual strong ignorability assumptions: 1) *unconfoundedness* $\{Y(0), Y(1)\} \perp\!\!\!\perp A|X$, 2) *strong overlap* ie. every patient has a strictly positive probability to receive each treatment, 3) *consistency*, and 4) *generalization* (detailed in B).

Estimating treatment effects with outcome models – g-computation [64] Should we know the two expected outcomes for a given X , we could compute the difference between them, which gives the causal effect of the treatment. These two expected outcomes can be estimated from observed data: the consistency 3 and unconfoundedness 1 assumptions imply the following equality:

$$\mathbb{E}_{Y(a) \sim \mathcal{D}^*}[Y(a)|X = x] = \mathbb{E}_{Y \sim \mathcal{D}}[Y|X = x, A = a] \quad (1)$$

On the left, the expectation is taken on the counterfactual unobserved distribution. On the right, the expectation is taken on the factual observed distribution conditionally on the treatment. For the rest of the paper, the expectations will always be taken on the factual observed distribution \mathcal{D} . This identification leads to outcome based estimators (ie. g-computation estimators [81]):

$$\begin{aligned} \tau &= \mathbb{E}_{Y \sim \mathcal{D}^*}[Y(1) - Y(0)|X = x] \\ &= \mathbb{E}_{Y \sim \mathcal{D}}[Y|A = 1] - \mathbb{E}_{Y \sim \mathcal{D}}[Y|A = 0] \end{aligned} \quad (2)$$

This equation builds on two quantities: the conditional expectancy of the outcome given the covariates and either treatment or no treatment, called *response function*:

$$\text{Response function} \quad \mu_a(x) \stackrel{\text{def}}{=} \mathbb{E}_{Y \sim \mathcal{D}}[Y|X = x, A = a]$$

Given a sample of data and the oracle response functions μ_0, μ_1 , the finite sum version of Equation 2 leads to an estimator of the ATE written:

$$\hat{\tau} = \frac{1}{n} \left(\sum_{i=1}^n \mu_1(x_i) - \mu_0(x_i) \right) \quad (3)$$

This estimator is an oracle **finite sum estimator** by opposition to the population expression of τ , $\mathbb{E}[\mu_1(x_i) - \mu_0(x_i)]$, which involves an expectation taken on the full distribution \mathcal{D} , which is observable but requires infinite data. For each estimator ℓ taking an expectation over \mathcal{D} , we use the symbol $\hat{\ell}$ to note its finite sum version.

Similarly to the ATE, at the individual level, the CATE:

$$\tau(x) = \mu_1(x) - \mu_0(x) \quad (4)$$

Robinson decomposition The *R-decomposition* of the outcome model plays an important role, [66]: introducing two quantities, the conditional mean outcome and the probability to be treated (known as propensity score [68]):

$$\text{Conditional mean outcome} \quad m(x) \stackrel{\text{def}}{=} \mathbb{E}_{Y \sim \mathcal{D}}[Y|X = x] \quad (5)$$

$$\text{Propensity score} \quad e(x) \stackrel{\text{def}}{=} \mathbb{P}[A = 1|X = x] \quad (6)$$

the outcome can be written

$$\begin{aligned} \text{R-decomposition} \quad y(a) &= m(x) + (a - e(x))\tau(x) + \varepsilon(x; a) \\ \text{with} \quad \mathbb{E}[\varepsilon(X; A)|X, A] &= 0 \end{aligned} \quad (7)$$

m and e are often called *nuisances* [15]; they are unknown.

Model-selection risks, oracle and feasible

Causal model selection We formalize model selection for causal estimation. Thanks to the g-formula identification (Equation 1), a given outcome model $f : \mathcal{X} \times \mathcal{A} \rightarrow \mathcal{Y}$ – learned from data or built from domain knowledge – induces feasible estimates of the ATE and CATE (eqs 3 and 4), $\hat{\tau}_f$ and $\hat{\tau}_f(x)$. Let $\mathcal{F} = \{f : \mathcal{X} \times \mathcal{A} \rightarrow \mathcal{Y}\}$ be a family of such estimators. Our goal is to select the best candidate in this family for the observed dataset O using a risk ℓ :

$$f_\ell^* = \underset{f \in \mathcal{F}}{\operatorname{argmin}} \ell(f, O) \quad (8)$$

We now detail possible risks ℓ , risks useful for causal model selection, and how to compute them.

The τ -risk: an oracle error risk As we would like to target the CATE, the following evaluation risk is natural (also called PEHE [72, 30]):

$$\tau\text{-risk}(f) \stackrel{\text{def}}{=} \mathbb{E}_{X \sim p(X)}[(\tau(X) - \hat{\tau}_f(X))^2] \quad (9)$$

Given observed data from $p(X)$, the expectation is computed with a finite sum, as in eq. 3, to give an estimated value $\widehat{\tau\text{-risk}}(f)$. However this risk is not feasible as the oracles $\tau(x)$ are not accessible with the observed data $(Y, X, A) \sim \mathcal{D}$.

Feasible error risks Table 1 lists *feasible* risks (Detailed in Appendix C), based on the prediction error of the outcome model and *observable* quantities. These observable, called nuisances are e – propensity score, eq 6 – and m – conditional mean outcome, eq 5. We give the definitions as *semi-oracles*, function of the true unknown nuisances, but later instantiate them with estimated nuisances, noted (\check{e}, \check{m}) . Semi-oracles risks are superscripted with the $*$ symbol.

Estimation and model selection procedure

Causal model selection (eq 8) may involve estimating various quantities from the observed data: the outcome model f , its

Table 1: Review of causal risks — The R -risk* is called τ -risk _{R} in Schuler et al. [73].

Risk	Equation	Reference
$mse(\tau(X), \tau_f(X)) = \tau\text{-risk}$	$\mathbb{E}_{X \sim p(X)}[(\tau(X) - \hat{\tau}_f(X))^2]$	Eq. 9 [30]
$mse(Y, f(X)) = \mu\text{-risk}$	$\mathbb{E}_{(Y, X, A) \sim \mathcal{D}}[(Y - f(X; A))^2]$	Def. 1 [73]
$\mu\text{-risk}_{IPW}^*$	$\mathbb{E}_{(Y, X, A) \sim \mathcal{D}} \left[\left(\frac{A}{e(X)} + \frac{1-A}{1-e(X)} \right) (Y - f(X; A))^2 \right]$	Def. 2 [41]
$\tau\text{-risk}_{IPW}^*$	$\mathbb{E}_{(Y, X, A) \sim \mathcal{D}} \left[\left(Y \left(\frac{A}{e(X)} - \frac{1-A}{1-e(X)} \right) - \hat{\tau}_f(X) \right)^2 \right]$	Def. 3 [89]
$U\text{-risk}^*$	$\mathbb{E}_{(Y, X, A) \sim \mathcal{D}} \left[\left(\frac{Y-m(X)}{A-e(X)} - \hat{\tau}_f(X) \right)^2 \right]$	Def. 4 [54]
$R\text{-risk}^*$	$\mathbb{E}_{(Y, X, A) \sim \mathcal{D}} \left[\left((Y - m(X)) - (A - e(X)) \hat{\tau}_f(X) \right)^2 \right]$	Def. 5 [54]

induced risk as introduced in the previous section, and possibly nuisances required by the risk. Given a dataset with N samples, we split out a train and a test sets $(\mathcal{T}, \mathcal{S})$. We fit each candidate estimator $f \in \mathcal{F}$ on \mathcal{T} . We also fit the nuisance models (\check{e}, \check{m}) on the train set \mathcal{T} , setting hyperparameters by a nested cross-validation before fitting the nuisance estimators with these parameters on the full train set. Causal quantities are then computed by applying the fitted candidates estimators $f \in \mathcal{F}$ on the test set \mathcal{S} . Finally, we compute the model-selection metrics for each candidate model on the test set. This procedure is described in Algorithm 1 and Figure 2.

Algorithm 1 Model selection procedure

Given train and test sets $(\mathcal{T}, \mathcal{S}) \sim \mathcal{D}$, a candidate estimators f , a causal metrics ℓ :

1. Prefit: Learn estimators for unknown nuisance quantities (\check{e}, \check{m}) on the training set \mathcal{T}
2. Fit: learn $\hat{f}(\cdot, a)$ on \mathcal{T}
3. Model selection: $\forall x \in \mathcal{S}$ predict $(\hat{f}(x, 1), \hat{f}(x, 0))$ and evaluate the estimator storing the metric value: $\ell(f, \mathcal{S})$ – possibly function of \check{e} and \check{m}




Figure 2: Estimation procedure for causal model selection.

The R-risk, a reweighted oracle metric

The R -risk can be rewritten as a rebalanced τ -risk.

This rewriting involves reweighted residuals: for each potential outcome, $a \in \{0; 1\}$, the variance conditionally on x is

[75]:

$$\sigma_y^2(x; a) \stackrel{\text{def}}{=} \int_y (y - \mu_a(x))^2 p(y | x = x; A = a) dy$$

Integrating over the population, we get the Bayes squared error: $\sigma_B^2(a) = \int_{\mathcal{X}} \sigma_y^2(x; a) p(x) dx$ and its propensity weighted version: $\tilde{\sigma}_B^2(a) = \int_{\mathcal{X}} \sigma_y^2(x; a) p(x; a) dx$. In case of a purely deterministic link between the covariates, the treatment, and the outcome, these residual terms are null.

Proposition 1 (R -risk as reweighted τ -risk) *Given an outcome model f , its R -risk appears as weighted version of its τ -risk (Proof in D.1):*

$$R\text{-risk}^*(f) = \int_x e(x)(1 - e(x)) (\tau(x) - \tau_f(x))^2 p(x) dx + \tilde{\sigma}_B^2(1) + \tilde{\sigma}_B^2(0) \quad (10)$$

The R -risk targets the oracle at the cost of an overlap reweighting and the addition of the reweighted Bayes residuals, which are independent of f . In good overlap regions the weights $e(x)(1 - e(x))$ are close to $\frac{1}{4}$, hence the R -risk is close to the desired gold-standard τ -risk. For randomized control trials, this weight is constant making the R -risk particularly suited for exploring heterogeneity (Appendix D.2)

Empirical Study

We evaluate the following causal metrics, oracle and feasible versions, presented in Table 1: $\widehat{\mu\text{-risk}}_{IPW}^*$, $\widehat{R\text{-risk}}^*$, $\widehat{U\text{-risk}}^*$, $\widehat{\tau\text{-risk}}_{IPW}^*$, $\widehat{\mu\text{-risk}}$, $\widehat{\mu\text{-risk}}_{IPW}$, $\widehat{R\text{-risk}}$, $\widehat{U\text{-risk}}$, $\widehat{\tau\text{-risk}}_{IPW}$. We benchmark the metrics in a variety of settings: many different simulated data generation processes and three semi-simulated datasets ¹.

Caussim: Extensive simulation settings

Data Generation We use simulated data, on which the ground-truth causal effect is known. Going beyond prior empirical studies of causal model selection [73, 1], we use many

¹ Scripts for the simulations and the selection procedure are available at <https://github.com/soda-inria/caussim>.

generative processes, to reach more general conclusions (as discussed in 21).

We generate the response functions using random bases extension, a common method in biostatistics, *eg* functional regression with splines [33, 58]. By allowing the function to vary at specific knots, we control the complexity of the non-linear outcome models. We use random approximation of Radial Basis Function (RBF) kernels [62] to generate the outcome and treatment functions. RBF use the same process as polynomial splines but replace polynomial by Gaussian kernels. Unlike polynomial, Gaussian kernels have decreasing influences in the input space. This avoids unrealistic divergences of the response surfaces at the ends of the feature space. We generate 1 000 datasets based on these functions, with random overlap parameters. Example shown in Figure 13 and details in F.1.

Family of candidate estimators We compare different candidate estimators that approximate imperfectly the data-generating process. To build such estimators, we first use a RBF expansion similar to the one used for data generation. We choose two random knots and apply a transformation of the raw data features with a Gaussian kernel. This step is referred as the featurization. Then, we fit a linear regression on these transformed features. We consider two ways of combining these steps for outcome mode; we use common nomenclature [40, 77] to refer to these different meta-learners that differ on how they model, jointly or not, the treated and the non treated:

- SLearner: A single learner for both population, taking the treatment as a supplementary covariate.
- SfLearner: A single set of basis functions is sampled at random for both populations, leading to a given feature space used to model both the treated and controls, then two separate different regressors are fitted on this shared representation.
- TLearner: Two completely different learners for each population, hence separate feature representations and regressors.

For the regression step, we fit a Ridge regression on the transformed features with 6 different choices of the regularization parameter $\lambda \in [10^{-3}, 10^{-2}, 10^{-1}, 1, 10^1, 10^2]$, coupled with a TLearner or a SfLearner. We sample 10 different random basis for learning and featurization yielding a family \mathcal{F} of 120 candidate estimators.

Semi-simulated datasets

Datasets We also use three benchmarks of the causal-inference literature, semi-simulated data adding a known synthetic causal effect to real –non synthetic– healthcare covariate. ACIC 2016 [20] is based on the Collaborative Perinatal Project [55], a RCT studying infants’ developmental disorders containing 4,802 individuals and 55 features. We used 770 dataset

instances: 10 random seeds for each of the 77 simulated settings for the treatment and outcomes. ACIC 2018 [79] simulated treatment and outcomes for the Linked Births and Infant Deaths Database (LBIDD) [47] with $D = 177$ covariates. We used all 432 datasets of size $N = 5\,000$. Twins [46] is an augmentation of real data on twin births and mortality rates [2]. There are $N = 11\,984$ samples, and $D = 50$ covariates for which we simulated 1,000 different treatment allocations. Datasets details are given in Appendix F.2.

Family of candidate estimators For these three datasets, the family of candidate estimators are gradient boosting trees for both the response surfaces and the treatment ² with S-learner, learning rate in $\{0.01, 0.1, 1\}$, and maximum number of leaf nodes in $\{25, 27, 30, 32, 35, 40\}$ resulting in a family of size 18.

Nuisance estimators Drawing inspiration from the TMLE literature that uses combination of flexible machine learning methods [74], we use a meta-learner as models for the nuisances \check{e} (respectively \check{m}): a stacked estimator of ridge and boosting classifiers (respectively regressions). We select hyper-parameters with randomized search on a validation set \mathcal{V} and keep them fix for model selection (F.3 lists hyperparameters).

Measuring overlap between treated and non treated

Good overlap between treated and control population is crucial for causal inference (positivity assumption 2). We introduce the Normalized Total Variation (NTV), a divergence based on the propensity score summarizing the overlap between both populations (Appendix E).

Results: factors driving good model selection

The R -risk is the best metric Figure 3 shows the agreement between the ideal ranking of outcome models given the oracle τ -risk and the different feasible causal metrics. We measure this agreement with relative³ Kendall tau κ (eq. 19) [38]. Given the importance of overlap in how well metrics approximate the oracle τ -risk, we separate strong and weak overlap.

Among all metrics, the classical mean squared error (ie. factual μ -risk) is worse and reweighting it with propensity score (μ -risk_{IPW}) does not bring much improvements. The R -risk, which includes a model of mean outcome and propensity scores, leads to the best performances. Interestingly, the U -risk, which uses the same nuisances, deteriorates in weak overlap, probably due to variance inflation when dividing by extreme propensity scores.

²Scikit-learn regressor, `HistGradientBoostingRegressor`, and classifier, `HistGradientBoostingClassifier`.

³To remove the variance across datasets (some datasets lead to easier model selection than others), we report values for one metric relative to the mean of all metrics for a given dataset instance: Relative $\kappa(\ell, \tau\text{-risk}) = \kappa(\ell, \tau\text{-risk}) - \text{mean}_\ell(\kappa(\ell, \tau\text{-risk}))$

Beyond rankings, the differences in terms of absolute ability to select the best model are large: The R-risk selects a model with a τ -risk only 1% higher than the best possible candidate for strong overlap on Caussim, but selecting with the μ -risk or μ -risk_{IPW}—as per machine-learning practice—leads to 10% excess risk and using τ -risk_{IPW}—as in some causal-inference methods [4, 27]—leads to 100% excess risk (Figure 16). Across datasets, the R -risk consistently decreases the risk compared to the μ -risk: 0.1% compared to 1% on ACIC2016, 1% compared to 20% on ACIC2018, and 0.05% compared to 1% on Twins.

Model selection is harder for low population overlap

Model selection for causal inference becomes more and more difficult with increasingly different treated and control populations (Figure 4). The absolute Kendall’s coefficient correlation with τ -risk drops from values around 0.9 (excellent agreement with oracle selection) to 0.6 on both Caussim and ACIC 2018 (F.4).

Nuisances can be estimated on the same data as outcome models Using the train set \mathcal{T} both to fit the candidate estimator and the nuisance estimates is a form of double dipping which can lead errors in nuisances correlated to that of outcome models [54]. In theory, these correlations can bias model selection and, strictly speaking, push to split out a third separated data set—a “nuisance set”—to fit the nuisance models. The drawback is that it depletes the data available for model estimation and selection. However, Figure 5 shows no substantial difference between a procedure with a separated nuisance set and the simpler shared nuisance-candidate set procedure.




Figure 3: **The R -risk is the best metric:** Relative Kendall’s τ agreement with τ -risk. Strong and Weak overlap correspond to the first and last tertiles of the overlap distribution measured with Normalized Total Variation eq. 16. F.4 presents the same results by adding semi-oracle risks in Figure 14, measured with absolute Kendall’s in Figure 15 and with τ -risk gains in Figure 16. Table 4 gives median and IQR of the relative Kendall.




Figure 4: **Model selection is harder for low population overlap:** Kendall’s τ agreement with τ -risk. Strong, medium and Weak overlap are the tertiles of the overlap measured with NTV eq. 16. F.4 presents results for all metrics in Figure 18 in absolute Kendall’s and continuous overlap values in Figure 15.




Figure 5: **Nuisances can be estimated on the same data as outcome models:** Results for the R -risk are similar between the **shared nuisances/candidate set** and the **separated nuisances set** procedures. Figure 17 details results for all metrics.

Stacked models are good overall estimators of nuisances For every risk, the oracle version recovers better the best estimator. However, stacked nuisances estimators (boosting and linear) lead to feasible metrics with close performances to the oracles ones: the corresponding estimators recover well-enough the true nuisances. One may wonder if simpler models for the nuisance could be useful, in particular in data-poor settings or when the true models are linear. Figure 6 compares causal model selection estimating nuisances with stacked estimators or linear model. It comprises the Twins data, where the true propensity model is linear, and a downsampled version of this data, to study a situation favorable to linear models. In these settings, stacked and linear estimations of the nuisances performs equivalently. Detailed analysis (Figure 20) confirms that using adaptive models—as built by stacking linear models and gradient-boosted trees—suffices to estimate nuisance.

Use 90% of the data to estimate outcome models, 10% to select them The analyst faces a compromise: given a finite data sample, should she allocate more data to estimate the outcome model, thus improving the quality of the outcome model but leaving little data for model selection. Or, she could choose




Figure 6: Stacked models are good overall estimators of the nuisances: Results are shown only for the R-risk; Figure 19 details every metrics. For Twins, where the true propensity model is linear, stacked and linear estimations of the nuisances performs equivalently, even for a downsampled version ($N=4,794$).

a bigger test set for model selection and effect estimation. For causal model selection, there is no established practice (as reviewed in G).

We investigate such tradeoff varying the ratio between train and test data size. For this, we first split out 30% of the data as a holdout set \mathcal{V} on which we use the oracle response functions to derive silver-standard estimates of causal quantities. We then use the standard estimation procedure on the remaining 70% of the data, splitting it into train \mathcal{T} and test \mathcal{S} of varying sizes. We finally measure the error between this estimate and the silver standard.

We consider two different analytic goals: estimating a average treatment effect –a single number used for policy making– and a CATE –a full model of the treatment effect as a function of covariates X . Given that the latter is a much more complex object than the former, the optimal train/test ratio might vary. To measure errors, we use for the ATE the relative absolute ATE bias between the ATE computed with the selected outcome model on the test set, and the true ATE as evaluated on the holdout set \mathcal{V} . For the CATE, we compare the τ -risk of the best selected model applied on the holdout set \mathcal{V} . We explore this trade-off for the ACIC 2016 dataset and the R-risk.

Figure 7 shows that a train/test ratio of 0.9/0.1 ($K=10$) or 0.8/0.2 ($K=5$) appears best to estimate CATE and ATE.

Discussion

Nuisance models: more gain than pain Predictive models are increasingly used to reason about causal effects, for instance in precision medicine to drive individualized decision. Our results highlight that they should be selected, validated, and tuned using different procedures and error measures than those classically used to assess prediction (estimating the so-called μ -risk). Rather, selecting the best outcome model according to the R -risk (eq. Definition 5) leads to more valid causal estimates. Estimating the R -risk requires a more complex procedure than standard cross-validation used e.g. in machine learning: it involves fitting nuisance models necessary for model evaluation, though our results show that these can be learned on the same set of data as the outcome model eval-




Figure 7: a) For CATE, a train/test ratio of 0.9/0.1 appears a good trade-off. b) For ATE, there is a small signal pointing also to 0.9/0.1 ($K=10$). for ATE. Experiences on 10 replications of all 78 instances of the ACIC 2016 data.

uated. The nuisance models must be well estimated (Figure 6). Our results show that using for nuisance models a flexible stacking-based family of estimator suffices for good model selection. For propensity score, we used the Brier score, which is minimized by the true individual probability. An easy mistake is to use calibration errors popular in machine learning [59, 93, 53, 48] as these select not for the individual posterior probability but for an aggregate error rate [57].

Extension to binary outcomes While we focused on continuous outcomes, in medicine, the target outcome is often a categorical variable such as mortality status or diagnosis. In this case, it may be interesting to focus on other estimands than the Average Treatment Effect $\mathbb{E}[Y(1)] - \mathbb{E}[Y(0)]$, for instance the relative risk $\frac{\mathbb{P}(Y(1)=1)}{\mathbb{P}(Y(0)=1)}$ or the odd ratio, $\frac{\mathbb{P}(Y(1)=1)/[1-\mathbb{P}(Y(1)=1)]}{\mathbb{P}(Y(0)=1)/[1-\mathbb{P}(Y(0)=1)]}$ are often used [6]. While the odds ratio is natural for case-control studies [69], other measures can reduce heterogeneity [16]. In the log domain, the ratios are written as a difference, the framework studied here can directly apply. In particular, the log odds ratio is estimated by the common cross-entropy loss (or log loss) as in logistic regression.

Conclusion

More R -risk to select models driving decisions Prediction models have flourished because their predictions can be easily demonstrated and validated on left-out data. But they require more careful validation for decision making, using a metric accounting for the putative intervention, the R -risk. Even when treated and untreated population differ little, as in RCTs, the R -risk brings a sizeable benefit. To facilitate better model se-

lection, we provide Python code⁴. Using the *R*-risk does make evaluation more complicated not only because the procedure is more involved, but also because each intervention requires a dedicated evaluation. However, such off-policy evaluation remains much less costly than the recommended good practice of impact evaluation testing the ability of a prediction model to actually guide patient health [28]. Also, the model-selection procedure puts no constraints on the models used to build predictive models: it opens the door to evaluating a wide range of models, from gradient boosting to convolutional neural, or language models.

Acknowledgments

We acknowledge fruitful discussions with Bénédicte Colnet.

Financial disclosure

None reported.

Conflict of interest

The authors declare no potential conflict of interests.

References

- [1] Alaa, Ahmed and Schaar, Mihaela Van Der. “Validating Causal Inference Models via Influence Functions”. In: *International Conference on Machine Learning* (May 24, 2019), pp. 191–201.
- [2] Almond, Douglas, Chay, Kenneth Y., and Lee, David S. “The Costs of Low Birth Weight”. In: *The Quarterly Journal of Economics* 120.3 (2005).
- [3] Altman, Douglas G, Vergouwe, Yvonne, Royston, Patrick, and Moons, Karel GM. “Prognosis and prognostic research: validating a prognostic model”. In: *Bmj* 338 (2009).
- [4] Athey, Susan and Imbens, Guido. “Recursive partitioning for heterogeneous causal effects”. In: *Proceedings of the National Academy of Sciences* 113.27 (2016), pp. 7353–7360.
- [5] Athey, Susan, Tibshirani, Julie, and Wager, Stefan. “Generalized random forests”. In: *Annals of Statistics* 47.2 (Apr. 2019), pp. 1148–1178.
- [6] Austin, Peter C and Stuart, Elizabeth A. “Estimating the effect of treatment on binary outcomes using full matching on the propensity score”. In: *Statistical methods in medical research* 26.6 (2017), pp. 2505–2525.
- [7] Austin, Peter C. “An Introduction to Propensity Score Methods for Reducing the Effects of Confounding in Observational Studies”. In: *Multivariate Behavioral Research* 3 (May 2011), pp. 399–424.
- [8] Austin, Peter C. and Stuart, Elizabeth A. “Moving towards best practice when using inverse probability of treatment weighting (IPTW) using the propensity score to estimate causal treatment effects in observational studies”. In: *Statistics in Medicine* 34.28 (Dec. 10, 2015), pp. 3661–3679.
- [9] Battocchi, Keith, Dillon, Eleanor, Hei, Maggie, Lewis, Greg, Oka, Paul, Oprescu, Miruna, and Syrgkanis, Vasilis. *EconML: A Python Package for ML-Based Heterogeneous Treatment Effects Estimation*. <https://github.com/microsoft/EconML>.
- [10] Beam, Andrew L and Kohane, Isaac S. “Big data and machine learning in health care”. In: *Jama* 319.13 (2018), pp. 1317–1318.
- [11] Blakely, Tony, Lynch, John, Simons, Koen, Bentley, Rebecca, and Rose, Sherri. “Reflection on modern methods: when worlds collide—prediction, machine learning and causal inference”. In: *International journal of epidemiology* 49.6 (2020), pp. 2058–2064.
- [12] Bouthillier, Xavier, Delaunay, Pierre, Bronzi, Mirko, Trofimov, Assya, Nichyporuk, Brennan, Szeto, Justin, Mohammadi Sepahvand, Nazanin, Raff, Edward, Madan, Kanika, Voleti, Vikram, Ebrahimi Kahou, Samira, Michalski, Vincent, Arbel, Tal, Pal, Chris, Varoquaux, Gael, and Vincent, Pascal. “Accounting for Variance in Machine Learning Benchmarks”. In: *Proceedings of Machine Learning and Systems* 3 (Mar. 15, 2021), pp. 747–769.
- [13] Casucci, Sabrina, Lin, Li, Hewner, Sharon, and Nikolaev, Alexander. “Estimating the causal effects of chronic disease combinations on 30-day hospital readmissions based on observational Medicaid data”. In: *Journal of the American Medical Informatics Association* 25.6 (2018), pp. 670–678.
- [14] Charlson, Mary E., Pompei, Peter, Ales, Kathy L., and MacKenzie, C.Ronald. “A new method of classifying prognostic comorbidity in longitudinal studies: Development and validation”. In: *Journal of Chronic Diseases* 40.5 (Jan. 1987), pp. 373–383.
- [15] Chernozhukov, Victor, Chetverikov, Denis, Demirer, Mert, Duflo, Esther, Hansen, Christian, Newey, Whitney, and Robins, James. “Double/Debiased Machine Learning for Treatment and Structural Parameters”. In: *The Econometrics Journal* (2018), p. 71.
- [16] Colnet, Bénédicte, Josse, Julie, Varoquaux, Gaël, and Scornet, Erwan. “Risk ratio, odds ratio, risk difference... Which causal measure is easier to generalize?” In: *arXiv preprint arXiv:2303.16008* (2023).
- [17] Curth, Alicia, Svensson, David, and Weatherall, James. “Really Doing Great at Estimating CATE? A Critical Look at ML Benchmarking Practices in Treatment Effect Estimation”. In: *Neurips Process 2021* (2021), p. 14.

⁴https://github.com/soda-inria/causal_model_selection

- [18] D'Amour, Alexander, Ding, Peng, Feller, Avi, Lei, Li-hua, and Sekhon, Jasjeet. "Overlap in observational studies with high-dimensional covariates". In: *Journal of Econometrics* 221.2 (2021), pp. 644–654.
- [19] Desai, Rishi J, Wang, Shirley V, Vaduganathan, Muthiah, Evers, Thomas, and Schneeweiss, Sebastian. "Comparison of machine learning methods with traditional models for use of administrative claims with electronic medical records to predict heart failure outcomes". In: *JAMA network open* 3.1 (2020), e1918962–e1918962.
- [20] Dorie, Vincent, Hill, Jennifer, Shalit, Uri, Scott, Marc, and Cervone, Dan. "Automated versus do-it-yourself methods for causal inference: Lessons learned from a data analysis competition". In: *Statistical Science* 34.1 (2019), pp. 43–68.
- [21] Fang, Gang, Annis, Izabela E, Elston-Lafata, Jennifer, and Cykert, Samuel. "Applying machine learning to predict real-world individual treatment effects: insights from a virtual patient cohort". In: *Journal of the American Medical Informatics Association* 26.10 (2019), pp. 977–988.
- [22] Fontana, Mark Alan, Lyman, Stephen, Sarker, Gourab K, Padgett, Douglas E, and MacLean, Catherine H. "Can machine learning algorithms predict which patients will achieve minimally clinically important differences from total joint arthroplasty?" In: *Clinical orthopaedics and related research* 477.6 (2019), p. 1267.
- [23] Gao, Zijun, Hastie, Trevor, and Tibshirani, Robert. "Assessment of heterogeneous treatment effect estimation accuracy via matching". In: *Statistics in Medicine* 17 (2021).
- [24] Gretton, Arthur, Borgwardt, Karsten M, Rasch, Malte J, Schölkopf, Bernhard, and Smola, Alexander. "A kernel two-sample test". In: *The Journal of Machine Learning Research* 13.1 (2012), pp. 723–773.
- [25] Grose, Elysia, Wilson, Samuel, Barkun, Jeffrey, Bertens, Kimberly, Martel, Guillaume, Balaa, Fady, and Khalil, Jad Abou. "Use of Propensity Score Methodology in Contemporary High-Impact Surgical Literature". In: *Journal of the American College of Surgeons* 230.1 (Jan. 1, 2020), 101–112.e2.
- [26] Gruber, Susan and van der Laan, Mark J. "tmle: An R Package for Targeted Maximum Likelihood Estimation". In: *Journal of Statistical Software* 51.13 (2012), pp. 1–35.
- [27] Gutierrez, Pierre and Gerardy, Jean-Yves. "Causal Inference and Uplift Modeling A review of the literature". In: *Proceedings of The 3rd International Conference on Predictive Applications and APIs*. Proceedings of Machine Learning Research 67 (2016), p. 14.
- [28] Hendriksen, Janneke MT, Geersing, Geert-Jan, Moons, Karel GM, and Groot, Joris AH de. "Diagnostic and prognostic prediction models". In: *Journal of Thrombosis and Haemostasis* 11 (2013), pp. 129–141.
- [29] Hernán, MA and Robins, JM. *Causal Inference: What If*. CRC Boca Raton, FL, 2020.
- [30] Hill, Jennifer L. "Bayesian Nonparametric Modeling for Causal Inference". In: *Journal of Computational and Graphical Statistics* 1 (Jan. 1, 2011), pp. 217–240.
- [31] Hoogland, Jeroen, IntHout, Joanna, Belias, Michail, Rovers, Maroeska M, Riley, Richard D, E. Harrell Jr, Frank, Moons, Karel GM, Debray, Thomas PA, and Reitsma, Johannes B. "A tutorial on individualized treatment effect prediction from randomized trials with a binary endpoint". In: *Statistics in medicine* 40.26 (2021), pp. 5961–5981.
- [32] Horng, Steven, Sontag, David A, Halpern, Yoni, Jernite, Yacine, Shapiro, Nathan I, and Nathanson, Larry A. "Creating an automated trigger for sepsis clinical decision support at emergency department triage using machine learning". In: *PloS one* 12.4 (2017), e0174708.
- [33] Howe, Chanelle J., Cole, Stephen R., Westreich, Daniel J., Greenland, Sander, Napravnik, Sonia, and Eron, Joseph J. "Splines for trend analysis and continuous confounder control". In: *Epidemiology (Cambridge, Mass.)* 22.6 (Nov. 2011), pp. 874–875.
- [34] Imbens, Guido W. and Rubin, Donald B. *Causal inference in statistics, social, and biomedical sciences*. Cambridge University Press, 2015.
- [35] Ionides, Edward L. "Truncated Importance Sampling". In: *Journal of Computational and Graphical Statistics* 17.2 (June 2008), pp. 295–311.
- [36] Jesson, Andrew, Mindermann, Sören, Shalit, Uri, and Gal, Yarin. "Identifying Causal-Effect Inference Failure with Uncertainty-Aware Models". In: *Advances in Neural Information Processing Systems* 33 (Oct. 22, 2020), pp. 11637–11649.
- [37] Johansson, Fredrik D, Shalit, Uri, Kallus, Nathan, and Sontag, David. "Generalization bounds and representation learning for estimation of potential outcomes and causal effects". In: *The Journal of Machine Learning Research* 23.1 (2022), pp. 7489–7538.
- [38] Kendall, M. G. "A new measure of rank correlation". In: *Biometrika* 30.1-2 (June 1938), pp. 81–93.
- [39] Khojaste-Sarakhs, M, Haghghi, Seyedhamidreza Shababi, Ghomi, SMT Fatemi, and Marchiori, Elena. "Deep learning for Alzheimer's disease diagnosis: A survey". In: *Artificial Intelligence in Medicine* (2022), p. 102332.
- [40] Künnel, Sören R., Sekhon, Jasjeet S., Bickel, Peter J., and Yu, Bin. "Metalearners for estimating heterogeneous treatment effects using machine learning". In: *Proceedings of the National Academy of Sciences* 116.10 (Mar. 5, 2019), pp. 4156–4165.

- [41] Laan, Mark J van der, Laan, MJ, and Robins, JM. *Unified methods for censored longitudinal data and causality*. Springer Science & Business Media, 2003.
- [42] Laan, Mark J. van der, Polley, Eric C., and Hubbard, Alan E. “Super Learner”. In: *Statistical Applications in Genetics and Molecular Biology* 6 (Sept. 16, 2007).
- [43] Laan, Mark J. van der and Rose, Sherri. *Targeted Learning*. Springer Series in Statistics. 2011.
- [44] Lamont, Andrea, Lyons, Michael D, Jaki, Thomas, Stuart, Elizabeth, Feaster, Daniel J, Tharmaratnam, Kukatharmini, Oberski, Daniel, Ishwaran, Hemant, Wilson, Dawn K, and Van Horn, M Lee. “Identification of predicted individual treatment effects in randomized clinical trials”. In: *Statistical methods in medical research* 27.1 (2018), pp. 142–157.
- [45] Loiseau, Nicolas, Trichelair, Paul, He, Maxime, Andreux, Mathieu, Zaslavskiy, Mikhail, Wainrib, Gilles, and Blum, Michael G. B. “External control arm analysis: an evaluation of propensity score approaches, G-computation, and doubly debiased machine learning”. In: *BMC Medical Research Methodology* 22 (2022).
- [46] Louizos, Christos, Shalit, Uri, Mooij, Joris, Sontag, David, Zemel, Richard, and Welling, Max. “Causal Effect Inference with Deep Latent-Variable Models”. In: *Advances in neural information processing systems* (Nov. 6, 2017).
- [47] MacDorman, M. F. and Atkinson, J. O. “Infant mortality statistics from the linked birth/infant death data set—1995 period data”. eng. In: *Monthly Vital Statistics Report* 46.6 Suppl 2 (Feb. 1998), pp. 1–22.
- [48] Minderer, Matthias, Djolonga, Josip, Romijnders, Rob, Hubis, Frances, Zhai, Xiaohua, Houlby, Neil, Tran, Dustin, and Lucic, Mario. “Revisiting the Calibration of Modern Neural Networks”. In: *Advances in Neural Information Processing Systems* 34 (2021), pp. 15682–15694.
- [49] Mooney, Stephen J and Pejaver, Vikas. “Big data in public health: terminology, machine learning, and privacy”. In: *Annual review of public health* 39 (2018), p. 95.
- [50] Naimi, Ashley I, Mishler, Alan E, and Kennedy, Edward H. “Challenges in Obtaining Valid Causal Effect Estimates with Machine Learning Algorithms”. In: *American Journal of Epidemiology* (2021).
- [51] Naimi, Ashley I and Whitcomb, Brian W. “Defining and Identifying Average Treatment Effects”. In: *American Journal of Epidemiology* (2023).
- [52] Nassif, Ali Bou, Talib, Manar Abu, Nasir, Qassim, Afadar, Yaman, and Elgendi, Omar. “Breast cancer detection using artificial intelligence techniques: A systematic literature review”. In: *Artificial Intelligence in Medicine* (2022), p. 102276.
- [53] Niculescu-Mizil, Alexandru and Caruana, Rich. “Predicting good probabilities with supervised learning”. en. In: *Proceedings of the 22nd international conference on Machine learning - ICML '05*. ACM Press, 2005, pp. 625–632.
- [54] Nie, Xinkun and Wager, Stefan. “Quasi-Oracle Estimation of Heterogeneous Treatment Effects”. In: *Biometrika* 108.2 (Dec. 13, 2017), pp. 299–319.
- [55] Niswander, Kenneth R. and Stroke, United States National Institute of Neurological Diseases and. *The Women and Their Pregnancies: The Collaborative Perinatal Study of the National Institute of Neurological Diseases and Stroke*. Google-Books-ID: A0bdVhlhDQkC. National Institute of Health, 1972. 562 pp.
- [56] Pedregosa, Fabian, Varoquaux, Gaël, Gramfort, Alexandre, Michel, Vincent, Thirion, Bertrand, Grisel, Olivier, Blondel, Mathieu, Prettenhofer, Peter, Weiss, Ron, Dubourg, Vincent, Vanderplas, Jake, Passos, Alexandre, Cournapeau, David, Brucher, Matthieu, Perrot, Matthieu, and Duchesnay, Édouard. “Scikit-learn: Machine Learning in Python”. In: *Journal of Machine Learning Research* 12.85 (2011), pp. 2825–2830.
- [57] Perez-Label, Alexandre, Morvan, Marine Le, and Varoquaux, Gaël. “Beyond calibration: estimating the grouping loss of modern neural networks”. In: *arXiv preprint arXiv:2210.16315* (2022).
- [58] Perperoglou, Aris, Sauerbrei, Willi, Abrahamowicz, Michal, and Schmid, Matthias. “A review of spline function procedures in R”. In: *BMC Medical Research Methodology* 19.1 (Mar. 6, 2019), p. 46.
- [59] Platt, John C. and Platt, John C. “Probabilistic Outputs for Support Vector Machines and Comparisons to Regularized Likelihood Methods”. In: *Advances in Large Margin Classifiers* (1999), pp. 61–74.
- [60] Poldrack, Russell A, Huckins, Grace, and Varoquaux, Gael. “Establishment of best practices for evidence for prediction: a review”. In: *JAMA psychiatry* 77.5 (2020), pp. 534–540.
- [61] Powers, Scott, Qian, Junyang, Jung, Kenneth, Schuler, Alejandro, Shah, Nigam H., Hastie, Trevor, and Tibshirani, Robert. “Some methods for heterogeneous treatment effect estimation in high dimensions”. In: *Statistics in Medicine* 37.11 (2018), pp. 1767–1787.
- [62] Rahimi, Ali and Recht, Benjamin. “Random Features for Large-Scale Kernel Machines”. In: *Advances in Neural Information Processing Systems*. Vol. 20. 2008.
- [63] Rajkomar, Alvin, Dean, Jeffrey, and Kohane, Isaac. “Machine learning in medicine”. In: *New England Journal of Medicine* 380.14 (2019), pp. 1347–1358.

- [64] Robins, James. “A new approach to causal inference in mortality studies with a sustained exposure period—application to control of the healthy worker survivor effect”. In: *Mathematical Modelling* 7.9 (Jan. 1, 1986), pp. 1393–1512.
- [65] Robins, James M. and Greenland, Sander. “The role of model selection in causal inference from non experimental data”. In: *American Journal of Epidemiology* 123.3 (Mar. 1986), pp. 392–402.
- [66] Robinson, P. M. “Root-N-Consistent Semiparametric Regression”. In: *Econometrica* 4 (1988), pp. 931–954.
- [67] Rolling, Craig A. and Yang, Yuhong. “Model selection for estimating treatment effects”. In: *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* 76.4 (Sept. 2014), pp. 749–769.
- [68] Rosenbaum, Paul R and Rubin, Donald B. “The central role of the propensity score in observational studies for causal effects”. In: *Biometrika* 70 (1983), pp. 41–55.
- [69] Rothman, KJ, Greenland, S, and Lash, TL. “Case-control studies, chapter 8”. In: *Modern epidemiology* (2008), pp. 111–127.
- [70] Rubin, Donald B. “Causal Inference Using Potential Outcomes”. In: *Journal of the American Statistical Association* 100.469 (Mar. 1, 2005), pp. 322–331.
- [71] Saito, Yuta and Yasui, Shota. “Counterfactual Cross-Validation: Stable Model Selection Procedure for Causal Inference Models”. In: *International Conference on Machine Learning*. PMLR. 2020, pp. 8398–8407.
- [72] Schulam, Peter and Saria, Suchi. “Reliable Decision Support using Counterfactual Models”. In: *Advances in neural information processing systems* 30 (Mar. 30, 2017).
- [73] Schuler, Alejandro, Baiocchi, Michael, Tibshirani, Robert, and Shah, Nigam. “A comparison of methods for model selection when estimating individual treatment effects”. In: *arXiv:1804.05146 [cs, stat]* (June 13, 2018).
- [74] Schuler, Megan S. and Rose, Sherri. “Targeted Maximum Likelihood Estimation for Causal Inference in Observational Studies”. In: *American Journal of Epidemiology* 185.1 (Jan. 1, 2017), pp. 65–73.
- [75] Shalit, Uri, Johansson, Fredrik D, and Sontag, David. “Estimating individual treatment effect: generalization bounds and algorithms”. In: *International Conference on Machine Learning*. PMLR. 2017, pp. 3076–3085.
- [76] Shen, Li, Margolies, Laurie R, Rothstein, Joseph H, Fluder, Eugene, McBride, Russell, and Sieh, Weiva. “Deep learning to improve breast cancer detection on screening mammography”. In: *Scientific reports* 9.1 (2019), p. 12495.
- [77] Shen, Lingjie, Geleijnse, Gijs, and Kaptein, Maurits. “RCTrep: An R Package for the Validation of Estimates of Average Treatment Effects”. In: *Journal of Statistical Software* (2023).
- [78] Shimoni, Yishai, Karavani, Ehud, Ravid, Sivan, Bak, Peter, Ng, Tan Hung, Alford, Sharon Hensley, Meade, Denise, and Goldschmidt, Yaara. “An Evaluation Toolkit to Guide Model Selection and Cohort Definition in Causal Inference”. In: *arXiv preprint arXiv:1906.00442* (2019).
- [79] Shimoni, Yishai, Yanover, Chen, Karavani, Ehud, and Goldschmidt, Yaara. “Benchmarking Framework for Performance-Evaluation of Causal Inference Analysis”. In: *arXiv:1802.05046 [cs, stat]* (Mar. 20, 2018).
- [80] Simon, Gregory E, Johnson, Eric, Lawrence, Jean M, Rossom, Rebecca C, Ahmedani, Brian, Lynch, Frances L, Beck, Arne, Waitzfelder, Beth, Ziebell, Rebecca, Penfold, Robert B, et al. “Predicting suicide attempts and suicide deaths following outpatient visits using electronic health records”. In: *American Journal of Psychiatry* 175.10 (2018), pp. 951–960.
- [81] Snowden, Jonathan M., Rose, Sherri, and Mortimer, Kathleen M. “Implementation of G-computation on a simulated data set: demonstration of a causal inference technique”. eng. In: *American Journal of Epidemiology* 173.7 (Apr. 2011), pp. 731–738.
- [82] Spasic, Irena, Nenadic, Goran, et al. “Clinical text data in machine learning: systematic review”. In: *JMIR medical informatics* 8.3 (2020), e17984.
- [83] Sperrin, Matthew, Jenkins, David, Martin, Glen P, and Peek, Niels. “Explicit causal reasoning is needed to prevent prognostic models being victims of their own success”. In: *Journal of the American Medical Informatics Association* 26.12 (2019), pp. 1675–1676.
- [84] Sriperumbudur, Bharath K., Fukumizu, Kenji, Gretton, Arthur, Schölkopf, Bernhard, and Lanckriet, Gert R. G. “On integral probability metrics, ϕ -divergences and binary classification”. In: *arXiv:0901.2698 [cs, math]* (Oct. 12, 2009).
- [85] Su, Xiaogang, Peña, Annette T, Liu, Lei, and Levine, Richard A. “Random forests of interaction trees for estimating individualized treatment effects in randomized trials”. In: *Statistics in medicine* 37.17 (2018), pp. 2547–2560.
- [86] Swaminathan, Adith and Joachims, Thorsten. “Counterfactual risk minimization: Learning from logged bandit feedback”. In: *International Conference on Machine Learning*. PMLR. 2015, pp. 814–823.
- [87] VanderWeele, Tyler J. “Principles of confounder selection”. In: *European journal of epidemiology* 34 (2019), pp. 211–219.
- [88] Varoquaux, Gaël and Colliot, Olivier. *Evaluating machine learning models and their diagnostic value*. 2022.
- [89] Wager, Stefan and Athey, Susan. “Estimation and Inference of Heterogeneous Treatment Effects using Random Forests”. In: *Journal of the American Statistical Association* 113.523 (July 3, 2018), pp. 1228–1242.

- [90] Wang, Hanyin, Li, Yikuan, Khan, Seema A, and Luo, Yuan. “Prediction of breast cancer distant recurrence using natural language processing and knowledge-guided convolutional neural network”. In: *Artificial intelligence in medicine* 110 (2020), p. 101977.
- [91] Wendling, T., Jung, K., Callahan, A., Schuler, A., Shah, N. H., and Gallego, B. “Comparing methods for estimation of heterogeneous treatment effects using observational data from health care databases”. In: *Statistics in Medicine* 23 (2018), pp. 3309–3324.
- [92] Yala, Adam, Lehman, Constance, Schuster, Tal, Portnoi, Tally, and Barzilay, Regina. “A deep learning mammography-based model for improved breast cancer risk prediction”. In: *Radiology* 292.1 (2019), pp. 60–66.
- [93] Zadrozny, Bianca and Elkan, Charles. “Obtaining calibrated probability estimates from decision trees and naive Bayesian classifiers”. In: (2001), p. 8.
- [94] Zhang, Zhenwei and Sejdić, Ervin. “Radiological images and machine learning: trends, perspectives, and prospects”. In: *Computers in biology and medicine* 108 (2019), pp. 354–370.




Figure 8: **Different outcome models lead to different estimation errors on the Average Treatment Effects**, on 77 classic simulations with known true causal effect [20]. The different models are ridge regression and random forests with different hyper-parameters (details A). The different configurations are plotted as a function of increasing difference between treated and untreated population –see section . There is no systematic best performer; data-driven model selection is important.

A Variability of ATE estimation on ACIC 2016

Figure 8 shows ATE estimations for six different models used in g-computation estimators on the 76 configurations of the ACIC 2016 dataset. Outcome models are fitted on half of the data and inference is done on the other half –ie. train/test with a split ratio of 0.5. For each configuration, and each model, this train test split was repeated ten times, yielding non parametric variance estimates [12]. Figure 8 shows large variations obtained across different outcome estimators on semi-synthetic datasets [20]. Flexible models such as random forests are doing well in most settings except when treated and untreated populations differ noticeably, in which case a linear model (ridge) is to be preferred. However random forests with different hyper-parameters (max depth= 2) yield poor estimates. A simple rule of thumb such as preferring flexible models does not work in general; model selection is needed.

Outcome models are implemented with `scikit-learn` [56] and the following hyper-parameters:

Outcome Model	Hyper-parameters grid
Random Forests	Max depth: [2, 10]
Ridge regression without treatment interaction	Ridge regularization: [0.1]
Ridge regression with treatment interaction	Ridge regularization: [0.1]

Table 2: Hyper-parameters grid used for ACIC 2016 ATE variability

B Causal assumptions

We assume the following four assumptions, referred as strong ignorability and necessary to assure identifiability of the causal estimands with observational data [70]:

Assumption 1 (Unconfoundedness)

$$\{Y(0), Y(1)\} \perp\!\!\!\perp A|X$$

This condition –also called ignorability– is equivalent to the conditional independence on $e(X)$ [68]: $\{Y(0), Y(1)\} \perp\!\!\!\perp A|e(X)$.

Assumption 2 (Overlap, also known as Positivity))

$$\eta < e(x) < 1 - \eta \quad \forall x \in \mathcal{X} \text{ and some } \eta > 0$$

The treatment is not perfectly predictable. Or with different words, every patient has a chance to be treated and not to be treated. For a given set of covariates, we need examples of both to recover the ATE.

As noted by [18], the choice of covariates X can be viewed as a trade-off between these two central assumptions. A bigger covariates set generally reinforces the ignorability assumption. In the contrary, overlap can be weakened by large \mathcal{X} because of the potential inclusion of instruments: variables only linked to the treatment which could lead to arbitrarily small propensity scores.

Assumption 3 (Consistency) The observed outcome is the potential outcome of the assigned treatment:

$$Y = AY(1) + (1 - A)Y(0)$$

Here, we assume that the intervention A has been well defined. This assumption focuses on the design of the experiment. It clearly states the link between the observed outcome and the potential outcomes through the intervention [29].

Assumption 4 (Generalization) The training data on which we build the estimator and the test data on which we make the estimation are drawn from the same distribution \mathcal{D}^* , also known as the “no covariate shift” assumption [36].

C Definitions of feasible risks

Definition 1 (Factual μ -risk) [75] This is the usual Mean Squared Error on the target y . It is what is typically meant by “generalization error” in supervised learning:

$$\mu\text{-risk}(f) = \mathbb{E}[(Y - f(X; A))^2]$$

Definition 2 (μ -risk $_{IPW}^*$) [41] Let the inverse propensity weighting function $w(x, a) = \frac{a}{e(x)} + \frac{1-a}{1-e(x)}$, we define the semi-oracle Inverse Propensity Weighting risk,

$$\mu\text{-risk}_{IPW}^*(f) = \mathbb{E}\left[\left(\frac{A}{e(X)} + \frac{1-A}{1-e(X)}\right)(Y - f(X; A))^2\right]$$

Definition 3 (τ -risk $_{IPW}^*$) [89] The CATE $\tau(x)$ can be estimated with a regression against inverse propensity weighted outcomes [4, 27, 89], the τ -risk $_{IPW}$.

$$\tau\text{-risk}_{IPW}^*(f) = \mathbb{E}\left[\left(Y \frac{A - e(X)}{e(X)(1 - e(X))} - \tau_f(X)\right)^2\right]$$

Definition 4 (U-risk *) [40, 54] Based on the Robinson decomposition –eq. 7, the U-learner uses the $A - e(X)$ term in the denominator. The derived risk is:

$$U\text{-risk}^*(f) = \mathbb{E}\left[\left(\frac{Y - m(X)}{A - e(X)} - \tau_f(X)\right)^2\right]$$

Note that extreme propensity weights in the denominator term might inflate errors in the numerator due to imperfect estimation of the mean outcome m .

Definition 5 (R-risk *) [54, 73] The R-risk also uses two nuisance m and e :

$$R\text{-risk}^*(f) = \mathbb{E}\left[\left((Y - m(X)) - (A - e(X))\tau_f(X)\right)^2\right]$$

It is also based on the Robinson decomposition –eq. 7.

D Proofs: Links between feasible and oracle risks

D.1 Reformulation of the R -risk as reweighted τ -risk

Proposition 1 (R -risk as reweighted τ -risk) **Proof 1** We consider the R -decomposition: [66],

$$y(a) = m(x) + (a - e(x))\tau(x) + \varepsilon(x; a) \quad (11)$$

Where $\mathbb{E}[\varepsilon(X; A)|X, A] = 0$ We can use it as plug in the R -risk formula:

$$\begin{aligned} R\text{-risk}(f) &= \int_{\mathcal{Y} \times \mathcal{X} \times \mathcal{A}} [(y - m(x)) - (a - e(x))\tau_f(x)]^2 p(y; x; a) dy dx da \\ &= \int_{\mathcal{Y} \times \mathcal{X} \times \mathcal{A}} [(a - e(x))\tau(x) + \varepsilon(x; a) - (a - e(x))\tau_f(x)]^2 p(y; x; a) dy dx da \\ &= \int_{\mathcal{X} \times \mathcal{A}} (a - e(x))^2 (\tau(x) - \tau_f(x))^2 p(x; a) dx da \\ &\quad + 2 \int_{\mathcal{Y} \times \mathcal{X} \times \mathcal{A}} (a - e(x))(\tau(x) - \tau_f(x)) \int_{\mathcal{Y}} \varepsilon(x; a) p(y | x; a) dy p(x; a) dx da \\ &\quad + \int_{\mathcal{X} \times \mathcal{A}} \int_{\mathcal{Y}} \varepsilon^2(x; a) p(y | x; a) dy p(x; a) dx da \end{aligned}$$

The first term can be decomposed on control and treated populations to force $e(x)$ to appear:

$$\begin{aligned} &\int_{\mathcal{X}} (\tau(x) - \tau_f(x))^2 \left[e(x)^2 p(x; 0) + (1 - e(x))^2 p(x; 1) \right] dx \\ &= \int_{\mathcal{X}} (\tau(x) - \tau_f(x))^2 \left[e(x)^2 (1 - e(x)) p(x) + (1 - e(x))^2 e(x) p(x) \right] dx \\ &= \int_{\mathcal{X}} (\tau(x) - \tau_f(x))^2 (1 - e(x)) e(x) [1 - e(x) + e(x)] p(x) dx \\ &= \int_{\mathcal{X}} (\tau(x) - \tau_f(x))^2 (1 - e(x)) e(x) p(x) dx. \end{aligned}$$

The second term is null since, $\mathbb{E}[\varepsilon(x, a)|X, A] = 0$.

The third term corresponds to the modulated residuals : $\tilde{\sigma}_B^2(0) + \tilde{\sigma}_B^2(1)$

D.2 Interesting special cases

Randomization special case If the treatment is randomized as in RCTs, $p(A = 1 | X = x) = p(A = 1) = p_A$, thus μ -risk_{IPW} takes a simpler form:

$$\mu\text{-risk}_{IPW} = \mathbb{E}_{(Y, X, A) \sim \mathcal{D}} \left[\left(\frac{A}{p_A} + \frac{1 - A}{1 - p_A} \right) (Y - f(X; A))^2 \right]$$

However, we still can have large differences between τ -risk and μ -risk_{IPW} coming from heterogeneous errors between populations as shown experimentally in Schuler et al. [73] and our results below.

Concerning the R -risk, replacing $e(x)$ by its randomized value p_A in Proposition 1 yields the oracle τ -risk up to multiplicative and additive constants:

$$R\text{-risk} = p_A (1 - p_A) \tau\text{-risk} + (1 - p_A) \sigma_B^2(0) + p_A \sigma_B^2(1)$$

Thus, selecting estimators with R -risk* in randomized setting controls the τ -risk. This explains the strong performances of R -risk in randomized setups [73] and is a strong argument to use it to estimate heterogeneity in RCTs.

Oracle Bayes predictor If we have access to the oracle Bayes predictor for the outcome ie. $f(x, a) = \mu(x, a)$, then all risks are equivalent up to the residual variance:

$$\tau\text{-risk}(\mu) = \mathbb{E}_{X \sim p(X)} [(\tau(X) - \tau_\mu(X))^2] = 0 \quad (12)$$

$$\begin{aligned}\mu\text{-risk}(\mu) &= \mathbb{E}_{(Y,X,A) \sim p(Y;X;A)}[(Y - \mu_A(X))^2] \\ &= \int_{\mathcal{X},\mathcal{A}} \varepsilon(x, a)^2 p(a|x) p(x) dx da \leq \sigma_B^2(0) + \sigma_B^2(1)\end{aligned}\tag{13}$$

$$\mu\text{-risk}_{IPW}(\mu) = \sigma_B^2(0) + \sigma_B^2(1) \quad \text{from Lemma ??}\tag{14}$$

$$R\text{-risk}(\mu) = \tilde{\sigma}_B^2(0) + \tilde{\sigma}_B^2(1) \leq \sigma_B^2(0) + \sigma_B^2(1)$$

from Proposition 1 (15)

Thus, differences between causal risks only matter in finite sample regimes. Universally consistent learners converge to the Bayes risk in asymptotic regimes, making all model selection risks equivalent. In practice however, choices must be made in non-asymptotic regimes.

E Measuring overlap

Motivation of the Normalized Total Variation Overlap is often assessed by comparing visually population distributions as in Figure 1 or computing standardized difference on each feature [7, 8]. While these methods are useful to decide if positivity holds, they do not yield a single measure. Rather, we compute the divergence between the population covariate distributions $\mathbb{P}(X|A = 0)$ and $\mathbb{P}(X|A = 1)$ [18, 37]. Computing overlap when working only on samples of the observed distribution, outside of simulation, requires a sophisticated estimator of discrepancy between distributions, as two data points never have the same exact set of features. Maximum Mean Discrepancy [24] is typically used in the context of causal inference [75, 37]. However it needs a kernel, typically Gaussian, to extrapolate across neighboring observations. We prefer avoiding the need to specify such a kernel, as it must be adapted to the data which is tricky with categorical or non-Gaussian features, a common situation for medical data.

For simulated and some semi-simulated data, we have access to the probability of treatment for each data point, which sample both densities in the same data point. Thus, we can directly use distribution discrepancy measures and rely on the Normalized Total Variation (NTV) distance to measure the overlap between the treated and control propensities. This is the empirical measure of the total variation distance [84] between the distributions, $TV(\mathbb{P}(X|A = 1), \mathbb{P}(X|A = 0))$. As we have both distribution sampled on the same points, we can rewrite it a sole function of the propensity score, a low dimensional score more tractable than the full distribution $\mathbb{P}(X|A)$:

$$\widehat{NTV}(e, 1 - e) = \frac{1}{2N} \sum_{i=1}^N \left| \frac{e(x_i)}{p_A} - \frac{1 - e(x_i)}{1 - p_A} \right| \tag{16}$$

Formally, we can rewrite NTV as the Total Variation distance between the two population distributions. For a population $O = (Y(A), X, A) \sim \mathcal{D}$:

$$\begin{aligned}NTV(O) &= \frac{1}{2N} \sum_{i=1}^N \left| \frac{e(x_i)}{p_A} - \frac{1 - e(x_i)}{1 - p_A} \right| \\ &= \frac{1}{2N} \sum_{i=1}^N \left| \frac{P(A = 1|X = x_i)}{p_A} - \frac{P(A = 0|X = x_i)}{1 - p_A} \right|\end{aligned}$$

Thus NTV approximates the following quantity in expectation over the data distribution \mathcal{D} :

$$\begin{aligned}NTV(\mathcal{D}) &= \int_{\mathcal{X}} \left| \frac{p(A = 1|X = x)}{p_A} - \frac{p(A = 0|X = x)}{1 - p_A} \right| p(x) dx \\ &= \int_{\mathcal{X}} \left| \frac{p(A = 1, X = x)}{p_A} - \frac{p(A = 0, X = x)}{1 - p_A} \right| dx \\ &= \int_{\mathcal{X}} |p(X = x|A = 1) - p(X = x|A = 0)| dx\end{aligned}$$

For countable sets, this expression corresponds to the Total Variation distance between treated and control populations covariate distributions : $TV(p_0(x), p_1(x))$.

Measuring overlap without the oracle propensity scores: For ACIC 2018, or for non-simulated data, the true propensity scores are not known. To measure overlap, we rely on flexible estimations of the Normalized Total Variation, using gradient boosting trees to approximate the propensity score. Empirical arguments for this plug-in approach is given in Figure 9.

Empirical arguments We show empirically that NTV is an appropriate measure of overlap by :

- Comparing the NTV distance with the MMD for Caussim which is gaussian distributed in Figure 11,
- Verifying that setups with penalized overlap from ACIC 2016 have a higher total variation distance than unpenalized setups in Figure 10.
- Verifying that the Inverse Propensity Weights extrema (the inverse of the ν overlap constant appearing in the overlap Assumption 2) positevely correlates with NTV for Caussim, ACIC 2016 and Twins in Figure 12. Even if the same value of the maximum IPW could lead to different values of NTV, we expect both measures to be correlated : the higher the extrem propensity weights, the higher the NTV.

Estimating NTV in practice Finally, we verify that approximating the NTV distance with a learned plug-in estimates of $e(x)$ is reasonable. We used either a logistic regression or a gradient boosting classifier to learn the propensity models for the three datasets where we have access to the ground truth propensity scores: Caussim, Twins and ACIC 2016. We respectively sampled 1000, 1000 and 770 instances of these datasets with different seeds and overlap settings. We first run a hyperparameter search with cross-validation on the train set, then select the best estimator. We refit on the train set this estimator with or without calibration by cross validation and finally estimate the normalized TV with the obtained model. This training procedure reflects the one described in Algorithm 1 where nuisance models are fitted only on the train set.

The hyper parameters are : learning rate $\in [1e-3, 1e-2, 1e-1, 1]$, minimum samples leaf $\in [2, 10, 50, 100, 200]$ for boosting and L2 regularization $\in [1e-3, 1e-2, 1e-1, 1]$ for logistic regression.

Results in Figure 9 comparing bias to the true normalized Total Variation of each dataset instances versus growing true NTV indicate that calibration of the propensity model is crucial to recover a good approximation of the NTV.

F Experiments

F.1 Details on the data generation process

We use Gaussian-distributed covariates and random basis expansion based on Radial Basis Function kernels. A random basis of RBF kernel enables modeling non-linear and complex relationships between covariates in a similar way to the well known spline expansion. The estimators of the response function are learned with a linear model on another random basis (which can be seen as a stochastic approximation of the full data kernel [62]). We carefully control the amount of overlap between treated and control populations, a crucial assumption for causal inference. Figure 13 illustrates 2D examples of the simulation.

- The raw features for both populations are drawn from a mixture of Gaussians: $\mathbb{P}(X) = p_A \mathbb{P}(X|A=1) + (1-p_A) \mathbb{P}(X|A=0)$ where $\mathbb{P}(x|A=a)$ is a rotated Gaussian:

$$\mathbb{P}(x|A=a) = W \cdot \mathcal{N}\left(\begin{bmatrix} (1-2a)\theta \\ 0 \end{bmatrix}; \begin{bmatrix} \sigma_0 & 0 \\ 0 & \sigma_1 \end{bmatrix}\right) \quad (17)$$

with θ a parameter controlling overlap (bigger yields poorer overlap), W a random rotation matrix and $\sigma_0^2 = 2$; $\sigma_1^2 = 5$.

This generation process allows to analytically compute the oracle propensity scores $e(x)$, to simply control for overlap with the parameter θ , the distance between the two Gaussian main axes and to visualize response surfaces.

- A basis expansion of the raw features increases the problem dimension. Using Radial Basis Function (RBF) Nystroem transformation ⁵, we expand the raw features into a transformed space. The basis expansion samples randomly a small number of representers in the raw data. Then, it computes an approximation of the full N-dimensional kernel with these basis components, yielding the transformed features $z(x)$. The number of basis functions –ie. *knots*–, controls the complexity of the ground-truth response surfaces and treatment. We first use this process to draw the non-treated response surface μ_0 and the causal effect τ . We then draw the observations from a mixture two Gaussians, for the treated and non treated. We vary the separation between the two Gaussians to control the overlap between treated and non-treated populations, an important parameter for causal inference (related to η in Proposition ??). Finally, we generate observed outcomes adding Gaussian noise.

⁵We use the Sklearn implementation, [56]




Figure 9: a) Without calibration, estimation of NTV is not trivial even for boosting models. b) Calibrated classifiers are able to recover the true Normalized Total Variation for all datasets where it is available.

Figure 10: NTV recovers well the overlap settings described in the ACIC paper [20]





Figure 11: Good correlation between overlap measured as normalized Total Variation and Maximum Mean Discrepancy (200 sampled Caussim datasets)






Figure 12: Maximal value of Inverse Propensity Weights increases exponentially with the overlap as measure by Normalized Total Variation.

More formally, we generate the basis following the original data distribution, $[b_1..b_D] \sim \mathbb{P}(x)$, with $D=2$ in our simulations. Then, we compute an approximation of the full kernel of the data generation process $RBF(x, \cdot)$ with $x \sim \mathbb{P}(x)$ with these representers: $z(x) = [RBF_\gamma(x, b_d)]_{d=1..D} \cdot Z^T \in \mathbb{R}^D$ with RBF_γ being the Gaussian kernel $K(x, y) = \exp(-\gamma \|x - y\|^2)$ and Z the normalization constant of the kernel basis, computed as the root inverse of the basis kernel $Z = [K(b_i, b_j)]_{i,j \in 1..D}^{-1/2}$

- Functions μ_0, τ are distinct linear functions of the transformed features:

$$\begin{aligned}\mu_0(x) &= [z(x); 1] \cdot \beta_\mu^T \\ \tau(x) &= [z(x); 1] \cdot \beta_\tau^T\end{aligned}$$


- Adding a Gaussian noise, $\varepsilon \sim \mathcal{N}(0, \sigma(x; a))$, we construct the potential outcomes: $y(a) = \mu_0(x) + a \tau(x) + \varepsilon(x, a)$

We generated 1000 instances of this dataset with uniformly random overlap parameters $\theta \in [0, 2.5]$.

F.2 Details on the semi-simulated datasets

ACIC 2018 [20]: The initial intervention was a child's birth weight ($A = 1$ if weight $< 2.5kg$), and outcome was the child's IQ after a follow-up period. The study contained $N = 4802$ data points with $D = 55$ features (5 binary, 27 count data,

Simulation: $D = 2$, $\theta = 0.7$, seed=8



One-dimensional cuts of the response surfaces




Figure 13: Example of the simulation setup in the input space with two knots –*i.e.* basis functions. The top panel shows the observations in feature space, while the bottom panel displays the two response surfaces on a 1D cut along the black lines drawn on the top panel.

and 23 continuous). They simulated 77 different setups varying parameters for treatment and response models, overlap, and interactions between treatment and covariates⁶. We used 10 different seeds for each setup, totaling 770 dataset instances.

ACIC 2018 [79]: Starting from data from the Linked Births and Infant Deaths Database (LBIDD) [47] with $D = 177$ covariates, treatment and outcome models are simulated with complex models to reflect different scenarios. The data do not provide the true propensity scores, so we evaluate only feasible metrics, which do not require this nuisance parameter. We used all 432 datasets⁷ of size $N = 5\,000$.

Twins [46]: It is an augmentation of real data on twin births and mortality rates [2]. There are $N = 11\,984$ samples (pairs of twins), and $D = 50$ covariates⁸. The outcome is the mortality and the treatment is the weight of the heavier twin at birth. This is a "true" counterfactual dataset [17] in the sense that we have both potential outcomes with each twin. They simulate the treatment with a sigmoid model based on GESTAT10 (number of gestation weeks before birth) and x the 45 other covariates:

$$\begin{aligned} \mathbf{t}_i \mid \mathbf{x}_i, \mathbf{z}_i &\sim \text{Bern}(\sigma(w_o^\top \mathbf{x} + w_h(\mathbf{z}/10 - 0.1))) \\ \text{with } w_o &\sim \mathcal{N}(0, 0.1 \cdot I), w_h \sim \mathcal{N}(5, 0.1) \end{aligned} \quad (18)$$

⁶Original R code available at <https://github.com/vdorie/aciccomp/tree/master/2016> to generate 77 simulations settings.

⁷Using the scaling part of the data, from github.com/IBM-HRL-MLHLS/IBM-Causal-Inference-Benchmarking-Framework

⁸We obtained the dataset from <https://github.com/AMLab-Amsterdam/CEVAE/tree/master/datasets/TWINS>

We add a non-constant slope in the sigmoid to control the overlap between treated and control populations. We sampled uniformly 1 000 different overlap parameters between 0 and 2.5, totaling 1 000 dataset instances. Unlike the previous datasets, only the overlap varies for these instances. The response surfaces are set by the original outcomes.

F.3 Model selection procedures

Nuisances estimation The nuisances are estimated with a stacked regressor inspired by the Super Learner framework, [42]). The hyper-parameters are optimized with a random search with following search grid detailed in Table 3. All implementations come from scikit-learn [56]. As extreme inverse propensity weights induce high variance, we use clipping [86, 35] to bound $\min(\hat{e}, 1 - \hat{e})$ away from 0 with a fixed $\eta = 10^{-10}$, ensuring strict overlap for numerical stability.

Model	Estimator	Hyper-parameters grid
Outcome, m	StackedRegressor (HistGradientBoostingRegressor, ridge)	ridge regularization: [0.0001, 0.001, 0.01, 0.1, 1, 10, 100] HistGradientBoostingRegressor learning rate: [0.01, 0.1, 1] HistGradientBoostingRegressor max leaf nodes: [10, 20, 30, 50]
Treatment, e	StackedClassifier (HistgradientBoostingClassifier, LogisticRegression)	LogisticRegression C: [0.0001, 0.001, 0.01, 0.1, 1, 10, 100] HistGradientBoostingClassifier learning rate: [0.01, 0.1, 1] HistGradientBoostingClassifier max leaf nodes: [10, 20, 30, 50]

Table 3: Hyper-parameters grid used for nuisance models

F.4 Additional Results

Definition of the Kendall's tau, κ The Kendall's tau is a widely used statistics to measure the rank correlation between two set of observations. It measures the number of concordant pairs minus the discordant pairs normalized by the total number of pairs. It takes values in the $[-1, 1]$ range.

$$\kappa = \frac{(\text{number of concordant pairs}) - (\text{number of discordant pairs})}{(\text{number of pairs})} \quad (19)$$

Values of relative $\kappa(\ell, \tau\text{-risk})$ compared to the mean over all metrics Kendall's as shown in the boxplots of Figure 3

Figure 14 - Results measured in relative Kendall's for feasible and semi-oracle risks Because of extreme propensity scores in the denominator and bayes error residuals in the numerator, the semi-oracle U -risk has poor performances at bad overlap. Estimating these propensity scores in the is feasible U -risk reduces the variance since clipping is performed.

Figure 15 - Results measured in absolute Kendall's

Figure 16 - Results measured as distance to the oracle tau-risk To see practical gain in term of τ -risk, we plot the results as the normalized distance between the estimator selected by the oracle τ -risk and the estimator selected by each causal metric.

Then, $\widehat{R}\text{-risk}^*$ is more efficient than all other metrics. The gain are substantial for every datasets.

Figure 17 - Stacked models for the nuisances is more efficient For each metrics the benefit of using a stacked model of linear and boosting estimators for nuisances compared to a linear model. The evaluation measure is Kendall's tau relative to the oracle R -risk* to have a stable reference between experiments. Thus, we do not include in this analysis the ACIC 2018 dataset since R -risk* is not available due to the lack of the true propensity score.

Figure 18 Low population overlap hinders model selection for all metrics

Figure 19 - Stacked models for the nuisances is more efficient For each metrics the benefit of using a stacked model of linear and boosting estimators for nuisances compared to a linear model. The evaluation measure is Kendall's tau relative to the oracle R -risk* to have a stable reference between experiments. Thus, we do not include in this analysis the ACIC 2018 dataset since R -risk* is not available due to the lack of the true propensity score.

Metric	Dataset	Strong Overlap		Weak Overlap	
		Median	IQR	Median	IQR
$\widehat{\mu\text{-risk}}$	Twins (N=11 984)	-0.32	0.12	-0.19	0.12
	ACIC 2016 (N=4 802)	-0.03	0.13	0.11	0.19
	Caussim (N=5 000)	-0.40	0.55	-0.16	0.31
	ACIC 2018 (N=5 000)	0.00	0.30	0.01	0.40
$\widehat{\mu\text{-risk}}_{IPW}$	Twins (N= 11 984)	-0.31	0.13	-0.17	0.12
	ACIC 2016 (N=4 802)	-0.02	0.13	0.11	0.19
	Caussim (N=5 000)	-0.34	0.50	0.09	0.31
	ACIC 2018 (N=5 000)	0.00	0.30	-0.01	0.43
$\widehat{\mu\text{-risk}}_{IPW}^*$	Twins (N= 11 984)	-0.32	0.13	-0.17	0.13
	ACIC 2016 (N=4 802)	-0.02	0.13	0.11	0.21
	Caussim (N=5 000)	-0.33	0.54	0.26	0.27
$\widehat{\tau\text{-risk}}_{IPW}$	Twins (N= 11 984)	0.13	0.12	0.27	0.12
	ACIC 2016 (N=4 802)	-0.07	0.18	0.05	0.31
	Caussim (N=5 000)	-0.19	0.43	-0.14	0.18
	ACIC 2018 (N=5 000)	-0.16	0.40	-0.11	0.66
$\widehat{\tau\text{-risk}}_{IPW}^*$	Twins (N= 11 984)	0.12	0.14	0.20	0.16
	ACIC 2016 (N=4 802)	-0.03	0.16	-0.09	0.43
	Caussim (N=5 000)	-0.15	0.46	-0.17	0.19
$\widehat{U\text{-risk}}$	Twins (N= 11 984)	0.13	0.12	0.02	0.25
	ACIC 2016 (N=4 802)	0.04	0.11	0.11	0.26
	Caussim (N=5 000)	0.04	0.43	-0.04	0.17
	ACIC 2018 (N=5 000)	0.12	0.26	-0.02	0.50
$\widehat{U\text{-risk}}^*$	Twins (N= 11 984)	0.25	0.08	-0.41	0.45
	ACIC 2016 (N=4 802)	0.08	0.13	-0.59	0.57
	Caussim (N=5 000)	0.46	0.12	0.02	0.44
$\widehat{R\text{-risk}}$	Twins (N= 11 984)	0.15	0.10	0.25	0.18
	ACIC 2016 (N=4 802)	0.07	0.12	0.22	0.15
	Caussim (N=5 000)	0.34	0.26	0.13	0.21
	ACIC 2018 (N=5 000)	0.13	0.27	0.21	0.47
$\widehat{R\text{-risk}}^*$	Twins (N= 11 984)	0.25	0.10	0.32	0.15
	ACIC 2016 (N=4 802)	0.12	0.12	0.25	0.15
	Caussim (N=5 000)	0.47	0.11	0.16	0.14

Table 4: Values of relative $\kappa(\ell, \tau\text{-risk})$ compared to the mean over all metrics Kendall's as shown in the boxplots of Figure 3

Figure 20 - Flexible models are performant in recovering nuisances even in linear setups

Selecting different seeds and parameters is crucial to draw conclusions One strength of our study is the various number of different simulated and semi-simulated datasets. We are convinced that the usual practice of using only a small number of generation processes does not allow to draw statistically significant conclusions.

Figure 21 illustrate the dependence of the results on the generation process for caussim simulations. We highlighted the different trajectories induced by three different seeds for data generation and three different treatment ratio instead of 1000 different seeds. The result curves are relatively stable from one setup to another for R -risk, but vary strongly for μ -risk and μ -risk_{IPW}.




Figure 14: The R -risk is the best metric: Relative Kendall's τ agreement with τ -risk. Strong and Weak overlap correspond to the first and last tertiles of the overlap distribution measured with Normalized Total Variation eq. 16.




Figure 15: Agreement with τ -risk ranking of methods function of overlap violation. The lines represent medians, estimated with a lowess. The transparent bands denote the 5% and 95% confidence intervals.






Figure 16: Metric performances by normalized tau-risk distance to the best method selected with τ -risk. All nuisances are learned with the same estimator stacking gradient boosting and ridge regression. Dotted and plain lines corresponds to 60% lowess quantile estimates. This choice of quantile allows to see better the oracle metrics lines for which outliers with a value of 0 distort the curves.


(a) Caussim



(b) ACIC 2016



(c) Twins



$$\kappa(\ell, \tau\text{-Risk}) - \kappa(R\text{-risk}^*, \tau\text{-Risk})$$

Figure 17: Results are similar between the [Shared nuisances/candidate set](#) and the [Separated nuisances set](#) procedure. The experience has not been run on the full metrics for Caussim due to computation costs.




Figure 18: **Low population overlap hinders causal model selection for all metrics:** Kendall’s τ agreement with τ -risk. Strong, medium and Weak overlap correspond to the tertiles of the overlap distribution measured with Normalized Total Variation eq. 16.




Figure 19: Learning the nuisances with [stacked models](#) (linear and gradient boosting) is important for successful model selection with R-risk. For Twins dataset, there is no improvement for [stacked models](#) compared to [linear models](#) because of the linearity of the propensity model.





Figure 20: **Flexible models are performant in recovering nuisances in the downsampled Twins dataset.** The propensity score is linear in this setup, making it particularly challenging for flexible models compared to linear methods.

Figure 21: Kendall correlation coefficients for each causal metric. Each (color, shape) pair indicates a different (treatment ratio, seed) of the generation process.



G Heterogeneity in practices for data split

Splitting the data is common when using machine learning for causal inference, but practices vary widely in terms of the fraction of data to allocate to train models, outcomes and nuisances, and to evaluate them.

Before even model selection, data splitting is often required for estimation of the treatment effect, ATE or CATE, for instance to compute the nuisances required to optimize the outcome model (as the R -risk, Definition 5). The most frequent choice is use 80% of the data to fit the models, and 20% to evaluate them. For instance, for CATE estimation, the R-learner has been introduced using K-folds with $K = 5$ and $K = 10$: 80% of the data (4 folds) to train the nuisances and the remaining fold to minimize the corresponding R-loss [54]. Yet, it has been implemented with $K=5$ in causallib [78] or $K=3$ in econML [9]. Likewise, for ATE estimation, Chernozhukov et al. [15] introduce doubly-robust machine learning, recommending $K=5$ based on an empirical comparison $K=2$. However, subsequent works use doubly robust ML with varying choices of K : Loiseau et al. [45] use $K=3$, Gao et al. [23] use $K=2$. In the econML implementation, K is set to 3 [9]. Naimi et al. [50] evaluate various machine-learning approaches –including R-learners– using $K=5$ and 10, drawing inspiration from the TMLE literature which sets $K=5$ in the TMLE package [26].

Causal model selection has been much less discussed. The only study that we are aware of, Schuler et al. [73], use a different data split: a 2-folds train/test procedure, training the nuisances on the first half of the data, and using the second half to estimate the R -risk and select the best treatment effect model.